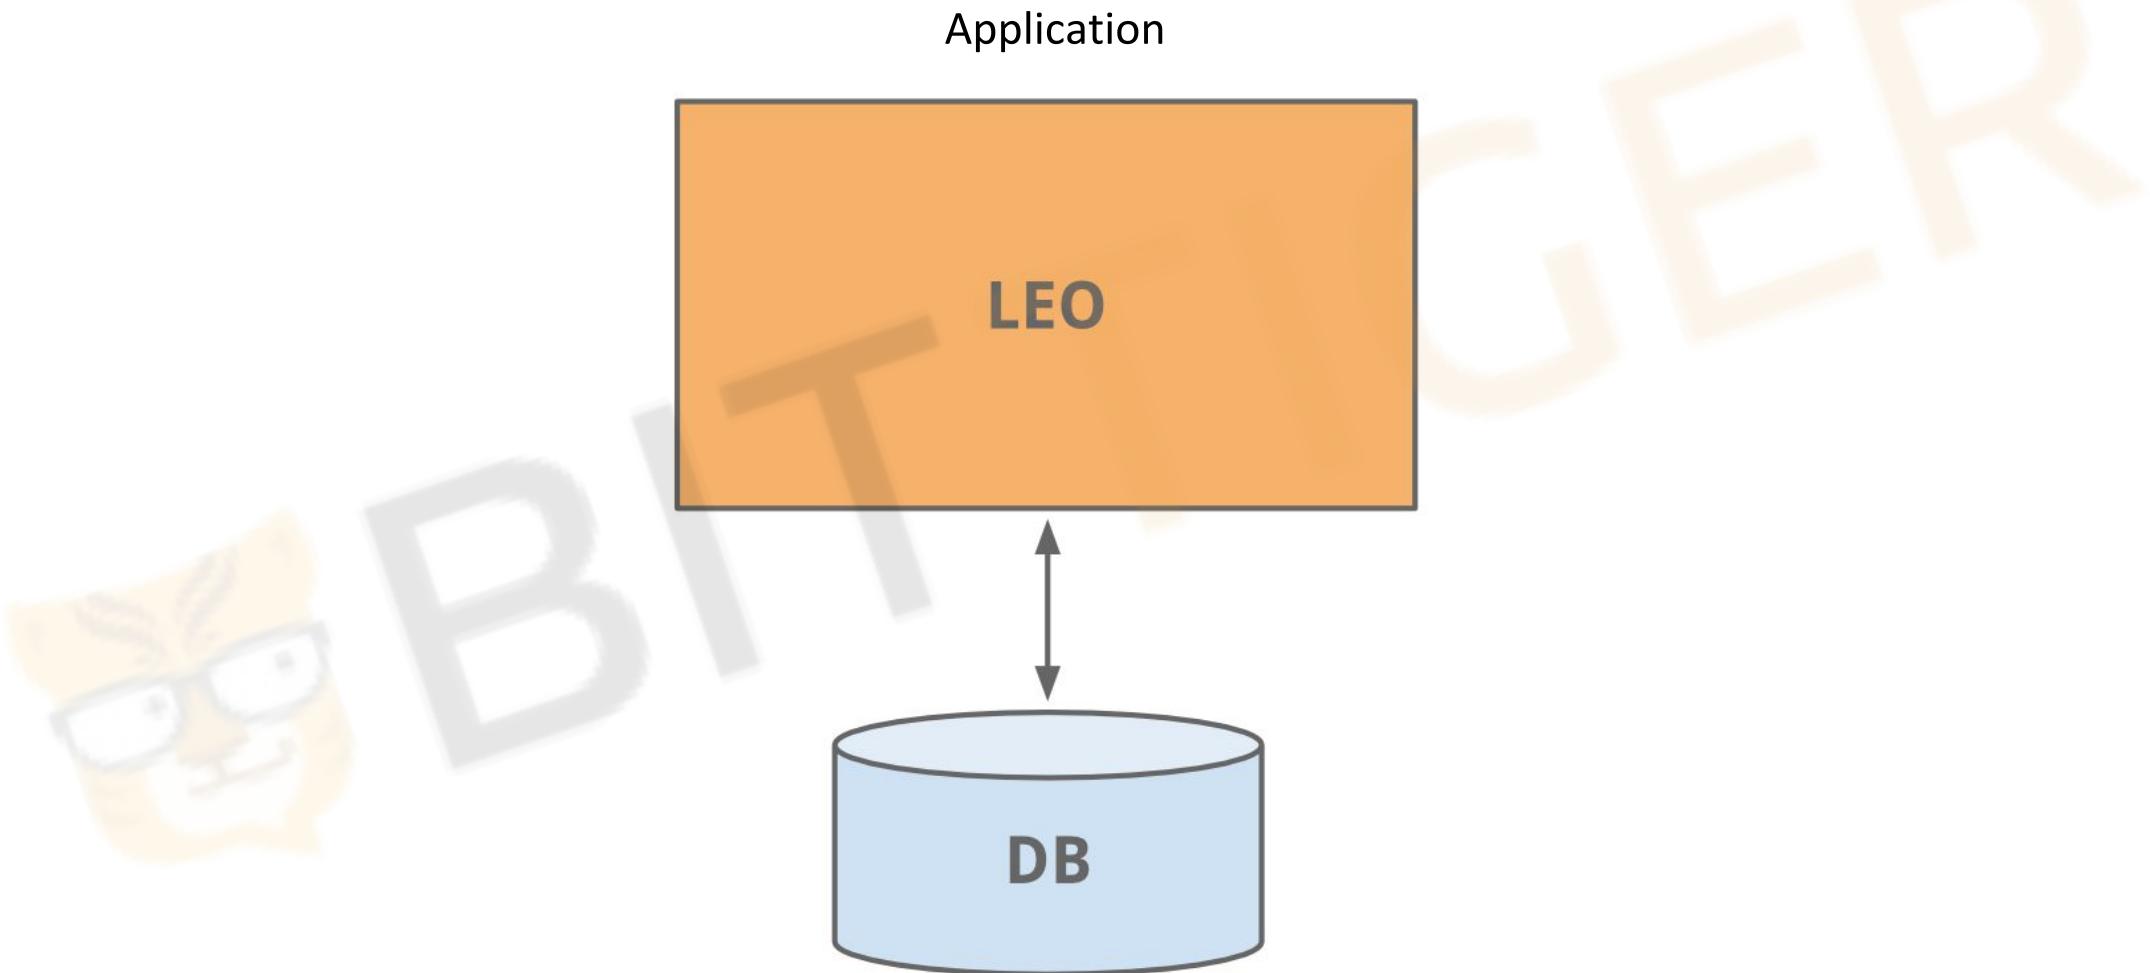


Design LinkedIn

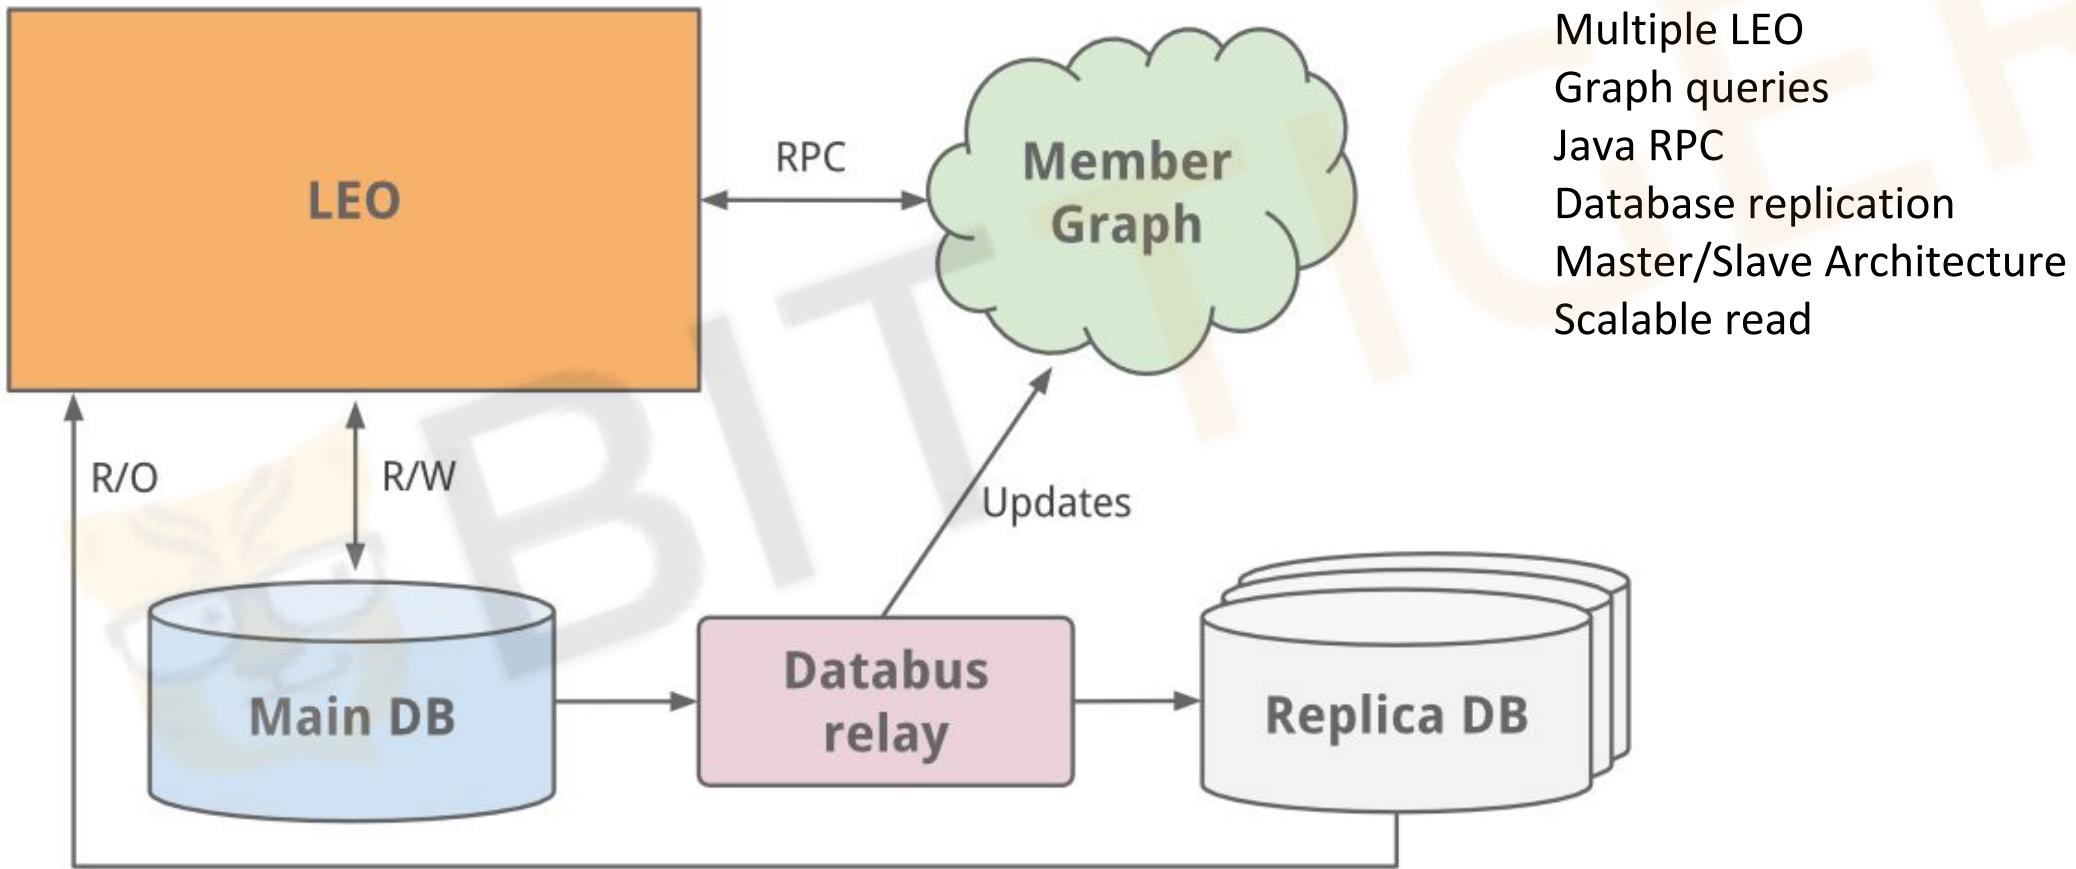
LinkedIn

- Profile
- Search
- Share
- Recommendation
- News feed
- Ads
- Messages

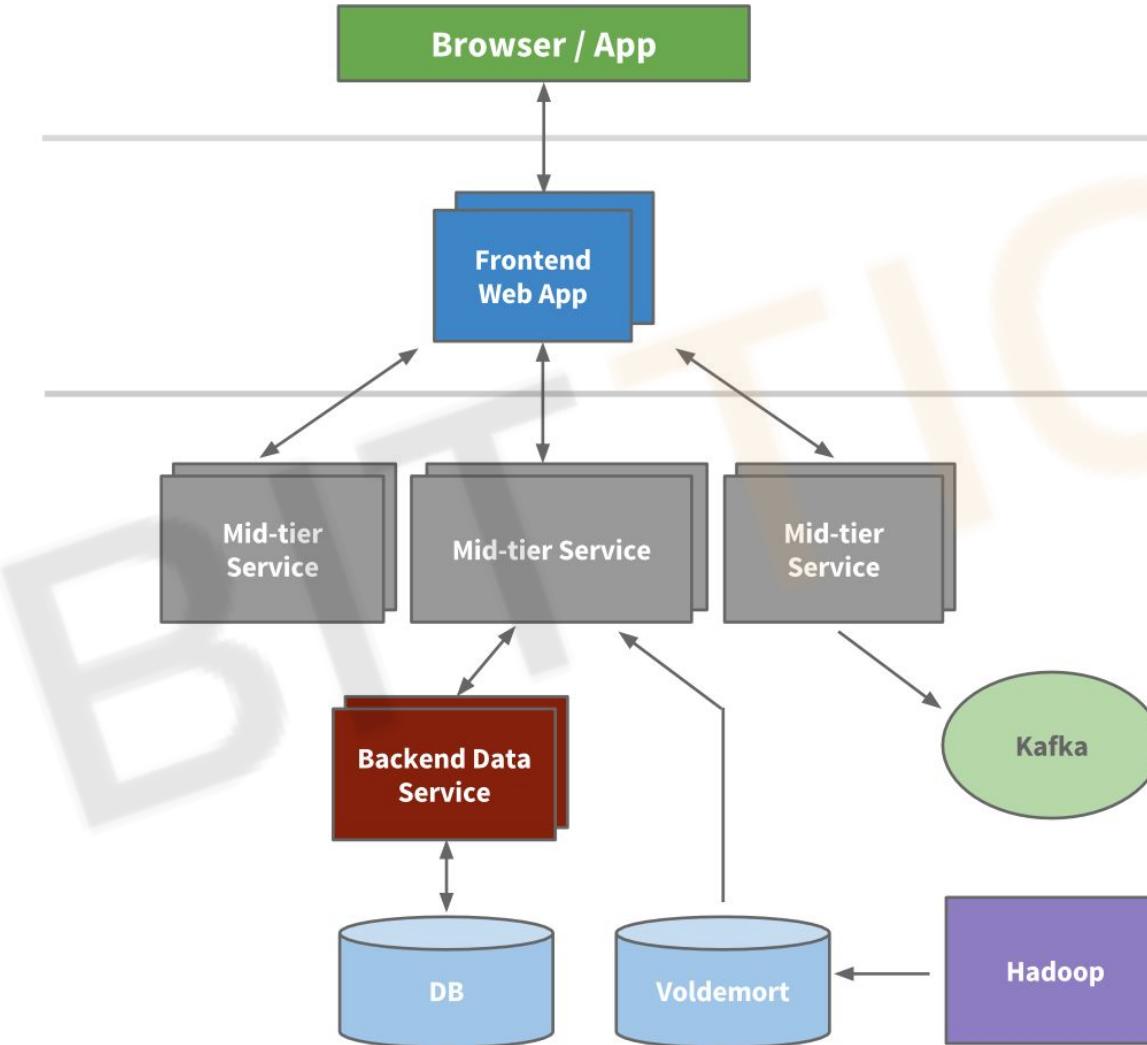
The early days



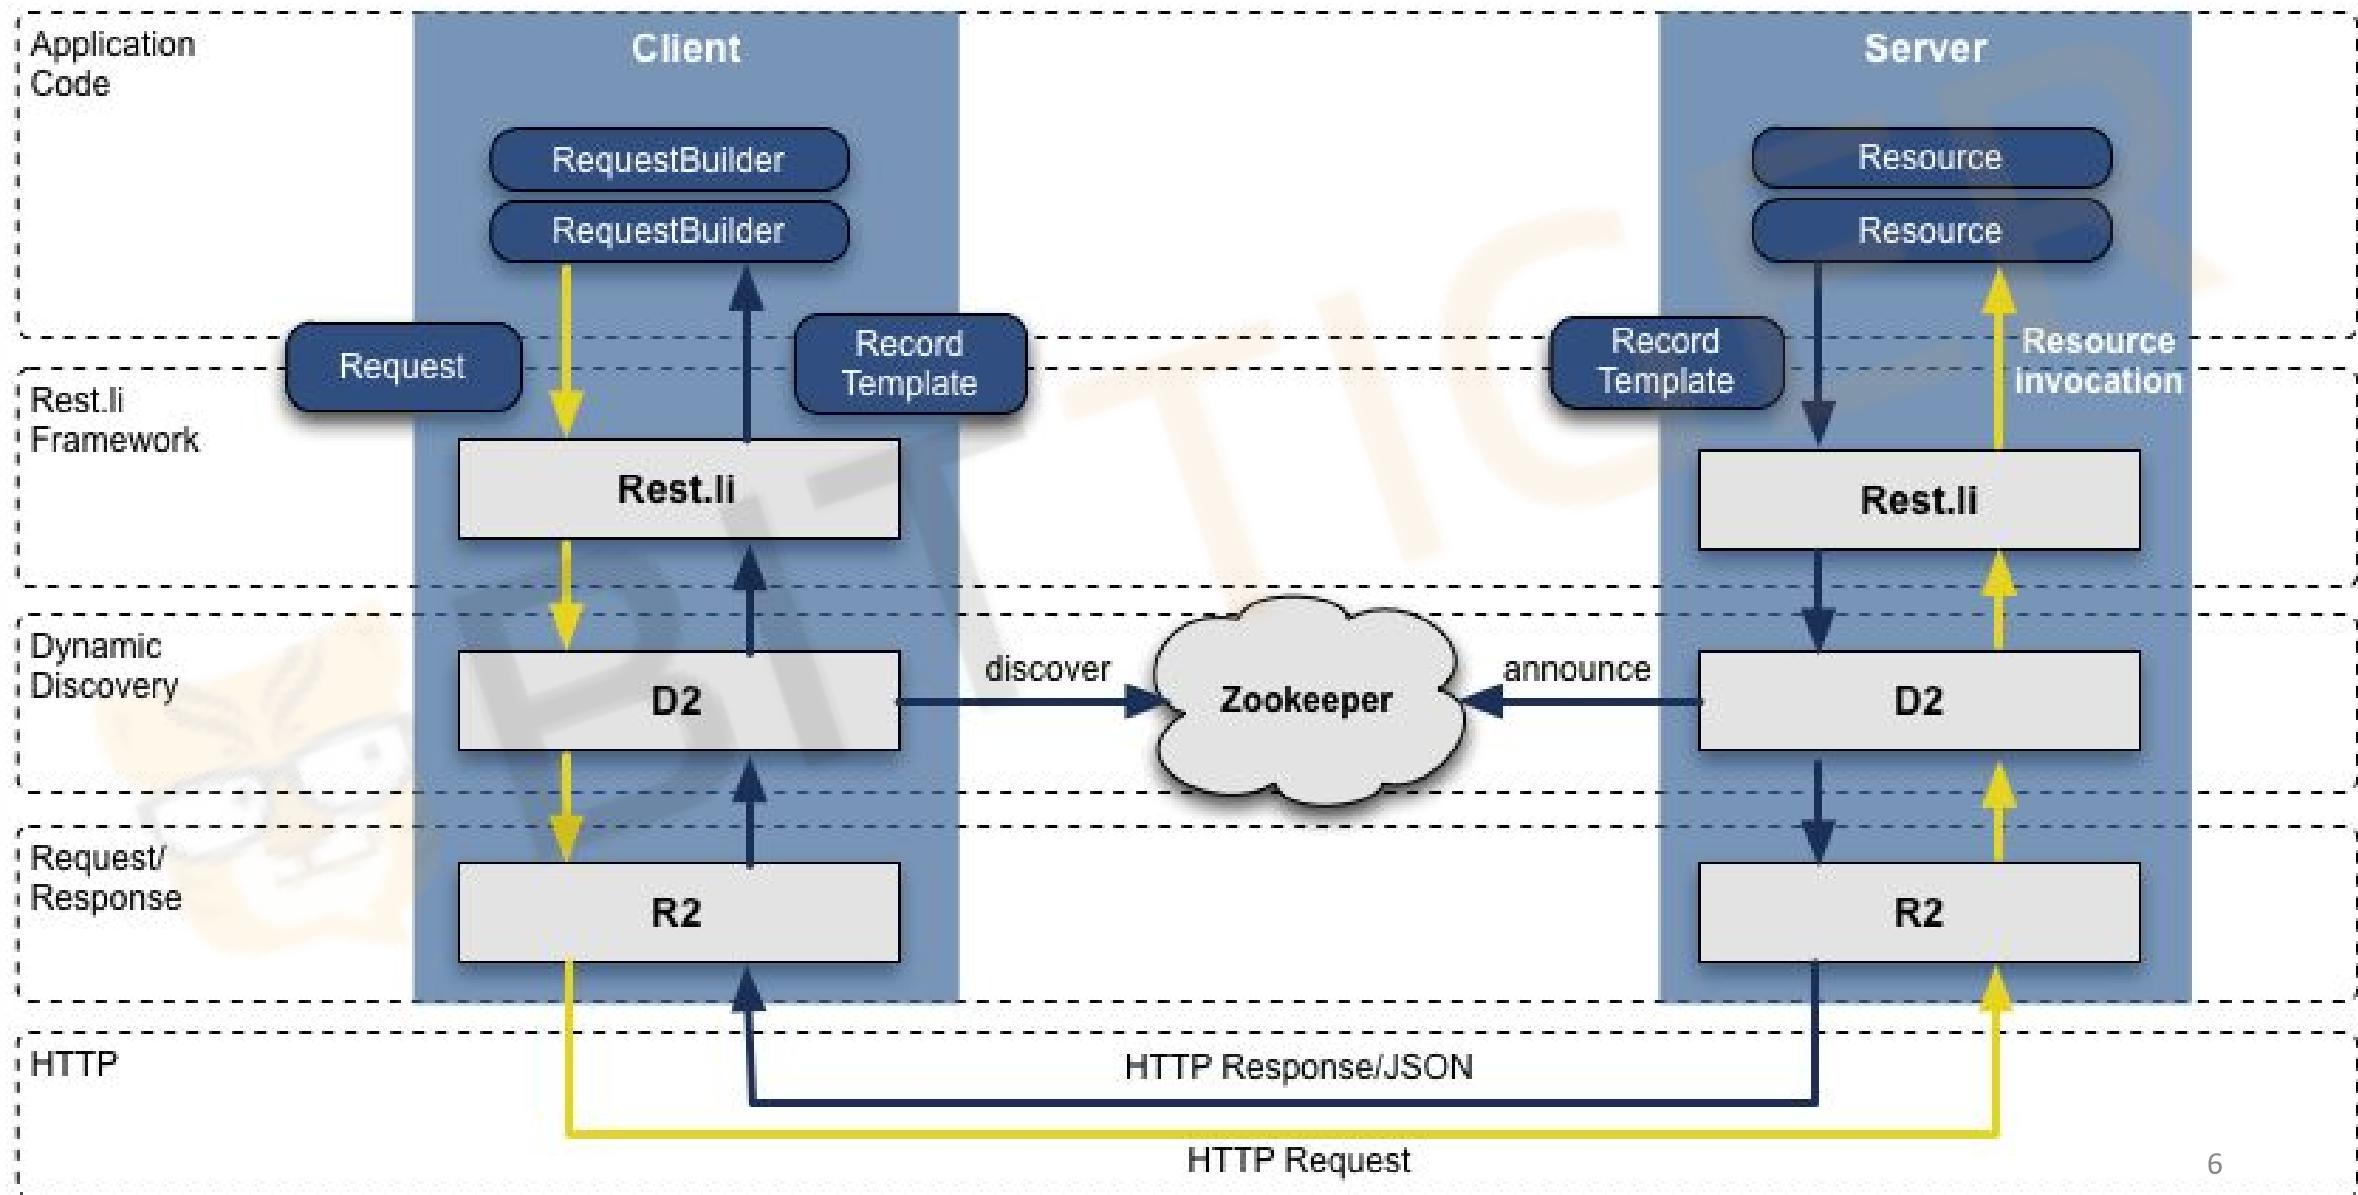
Distributed systems



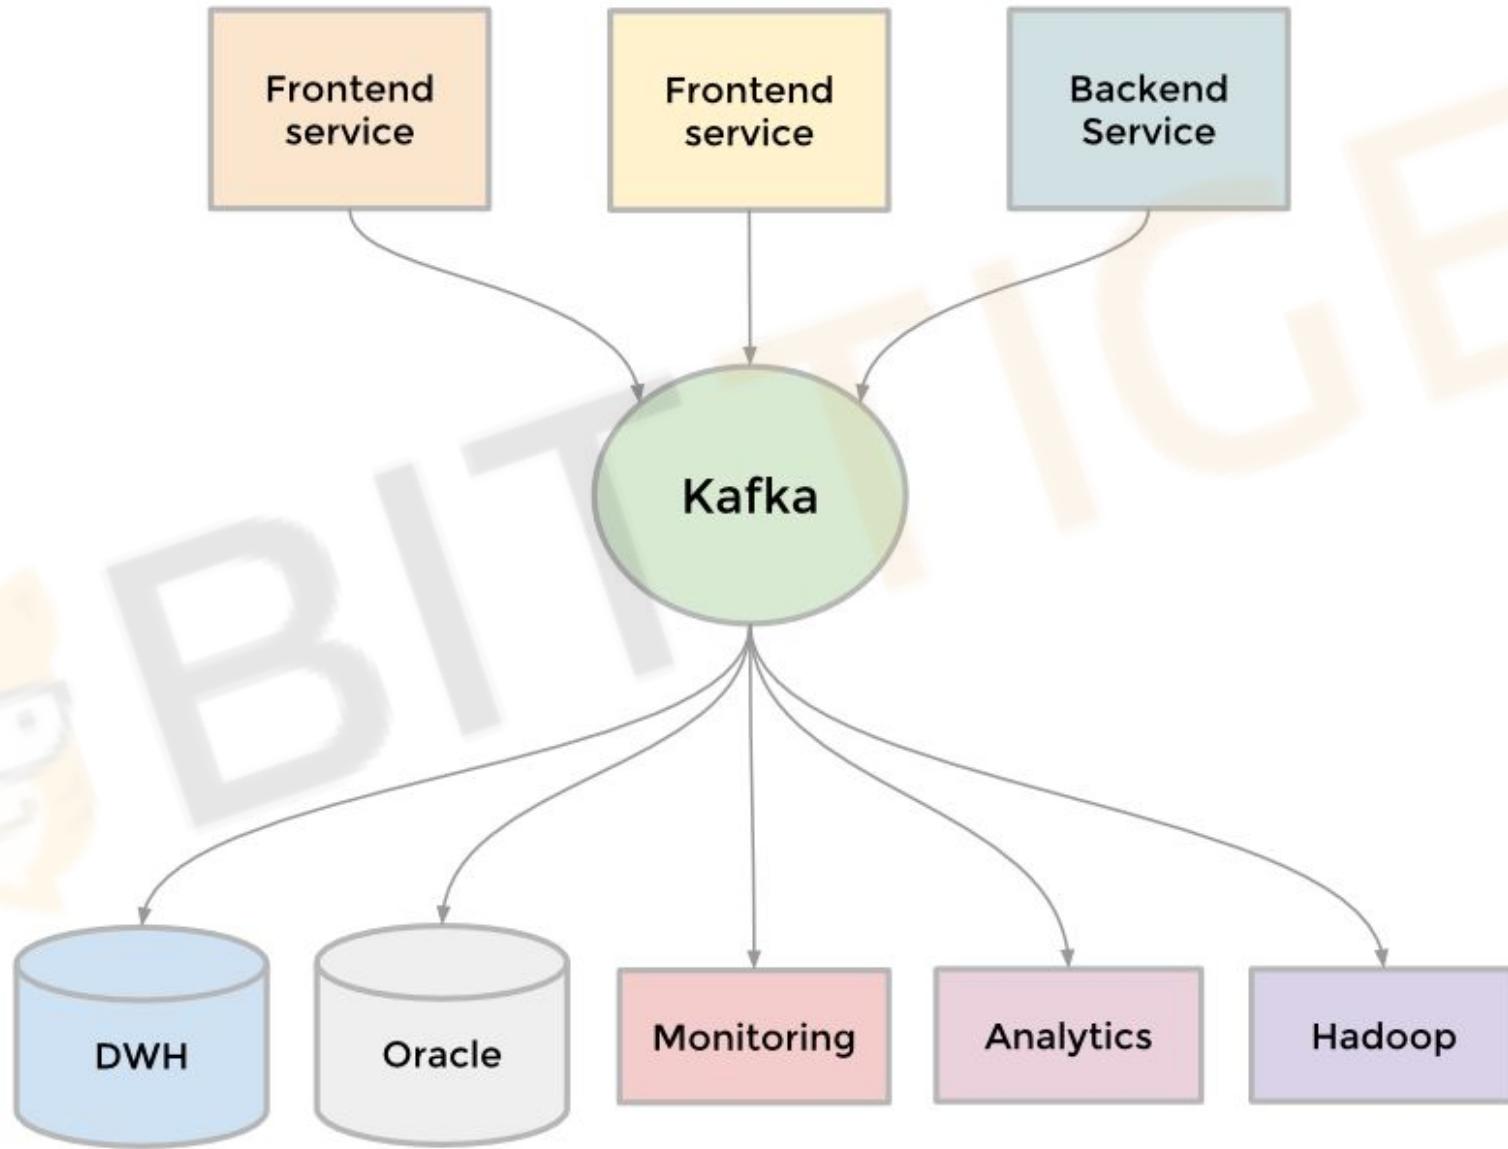
Service oriented architecture



Standardize REST API



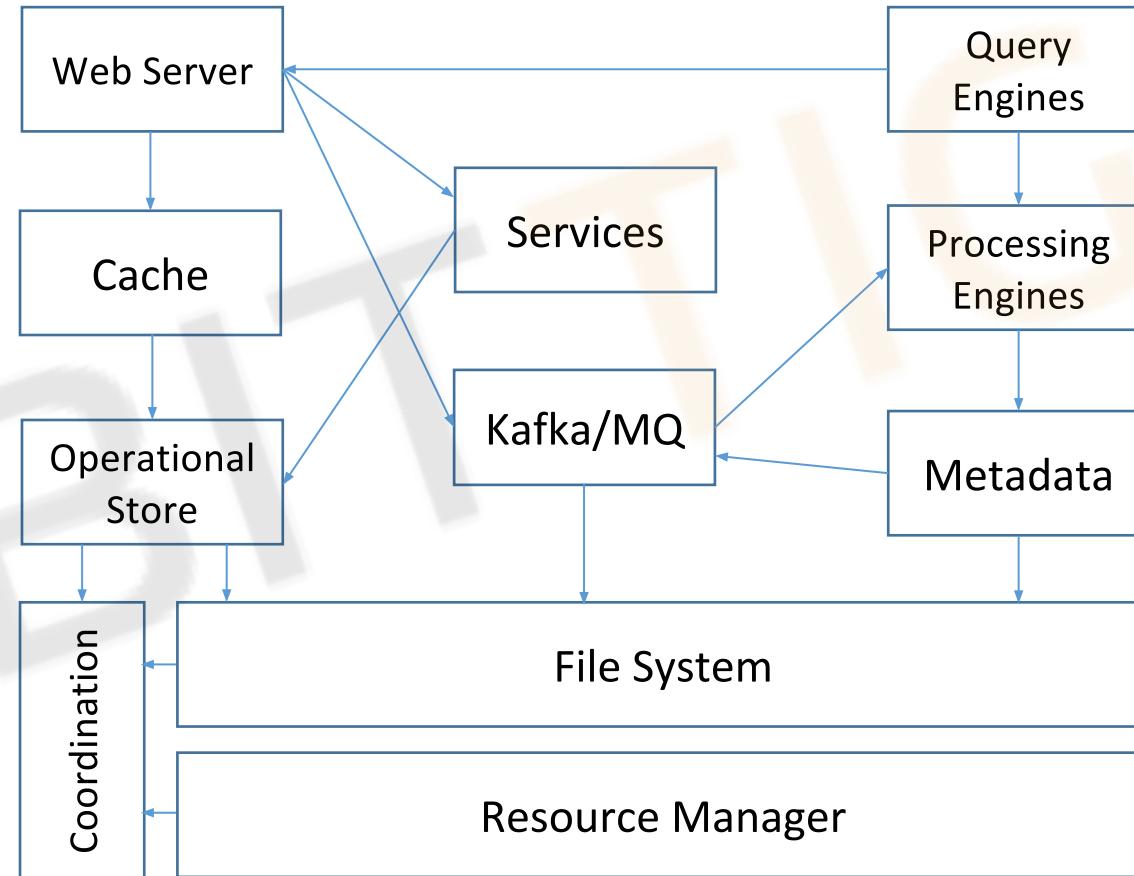
Data pipeline



Scaling LinkedIn

- Distributed systems
 - member graph
- Separation of concern
 - service oriented architecture
- Optimization of common cases
 - read/write separation
- Asynchronous processing
 - Kafka
- Standard
 - Rest.li

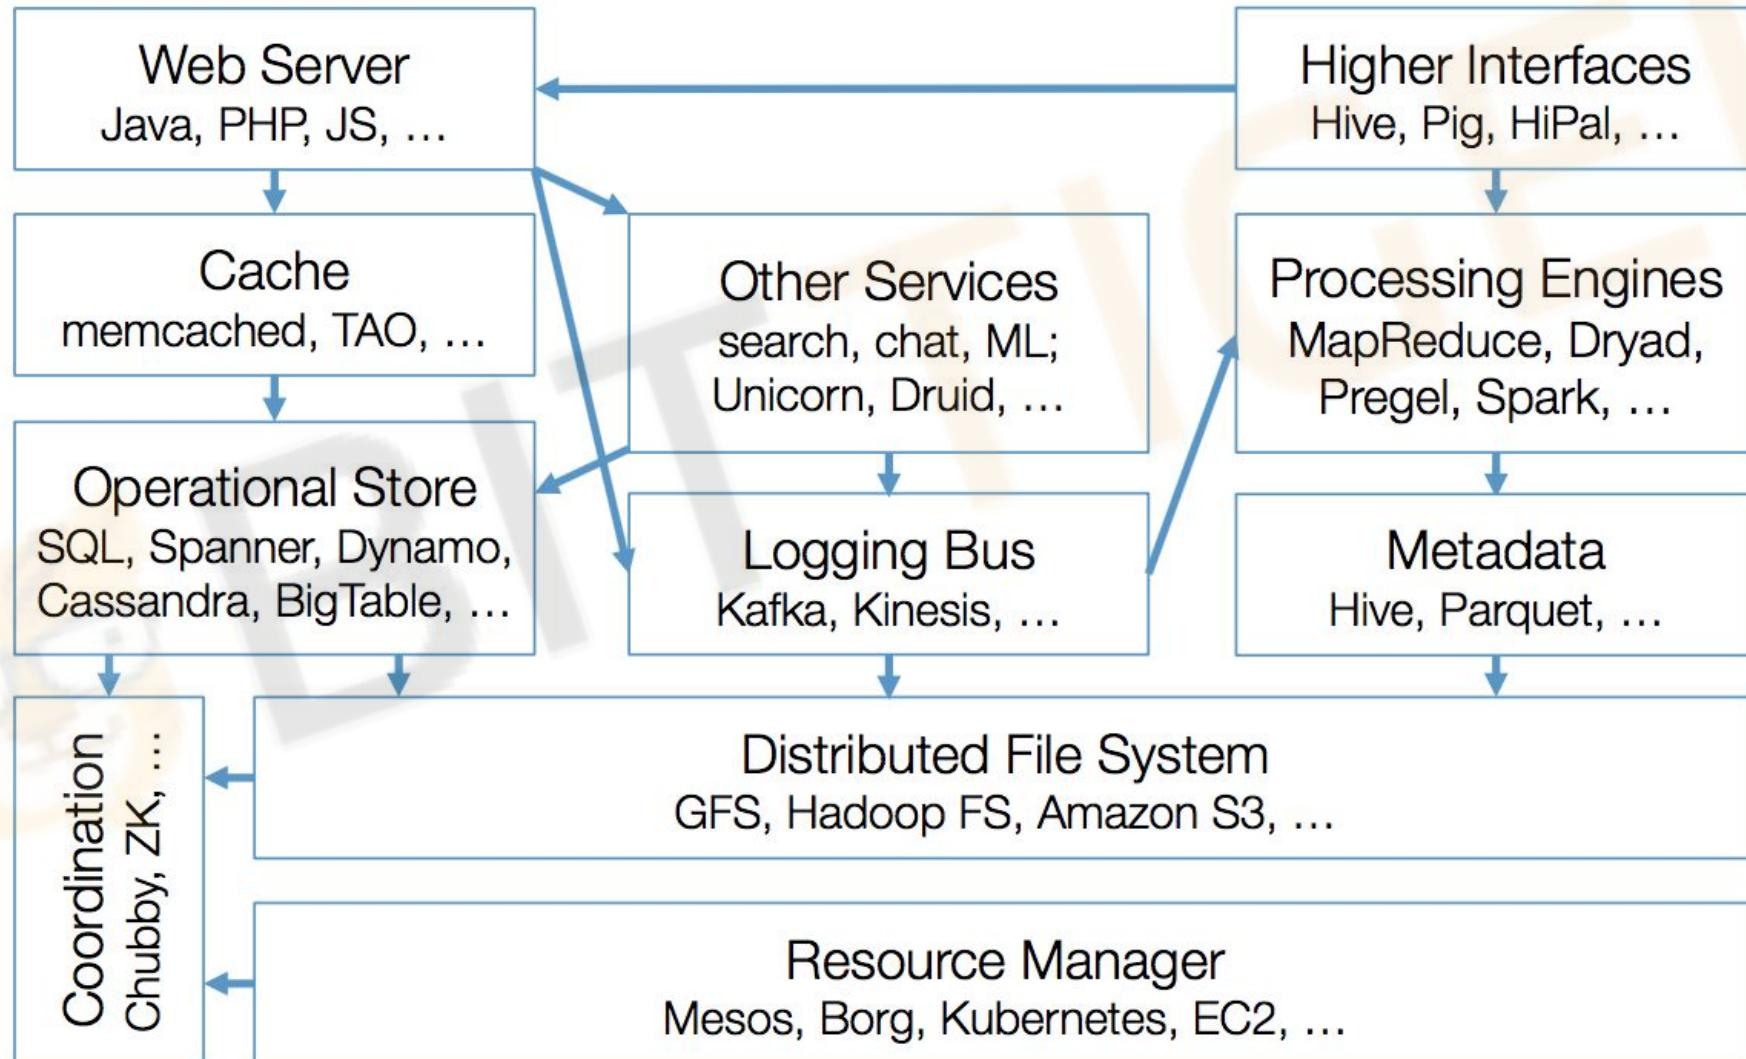
Scalable architecture



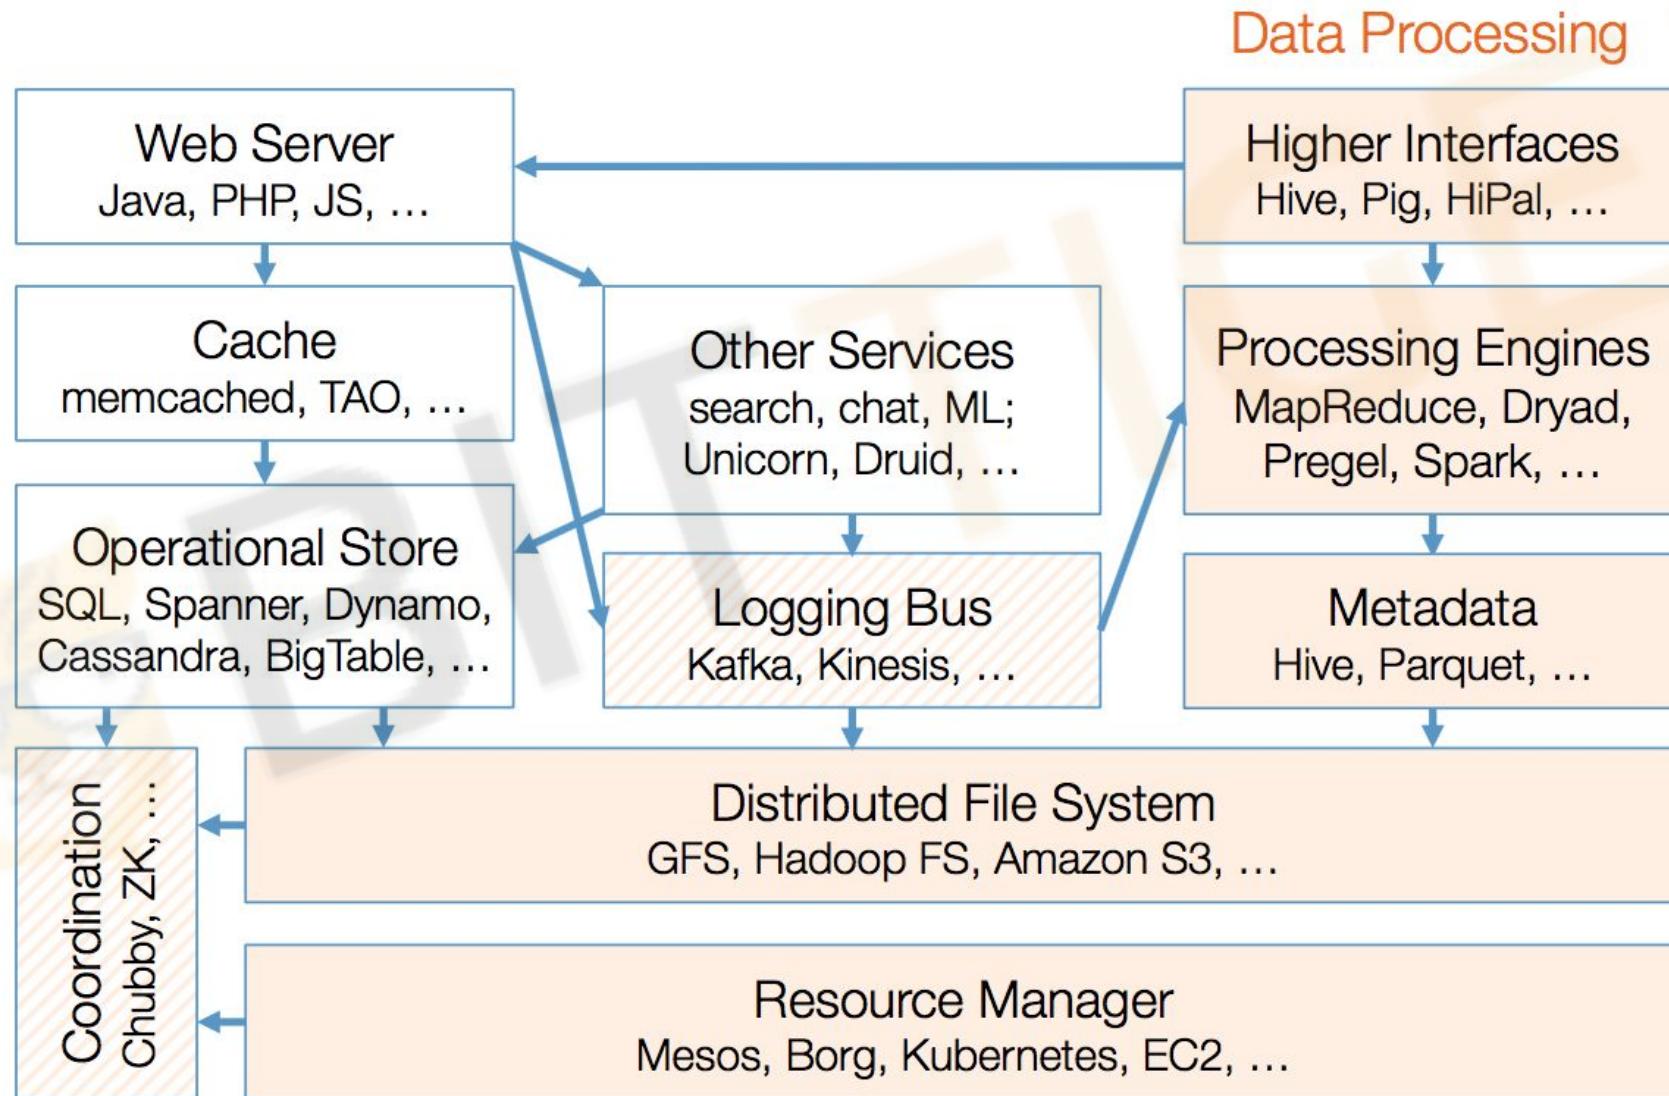
LinkedIn software stack

- Operational Store
 - Espresso
 - Vodemort
- Services
 - Search: Galene
 - Graph: member graph service
 - SOA: Rest.li
- Processing engines
 - Offline: Hadoop, Spark
 - Online: Samza
- Metadata
 - Schema Registry
 - Hive metastore

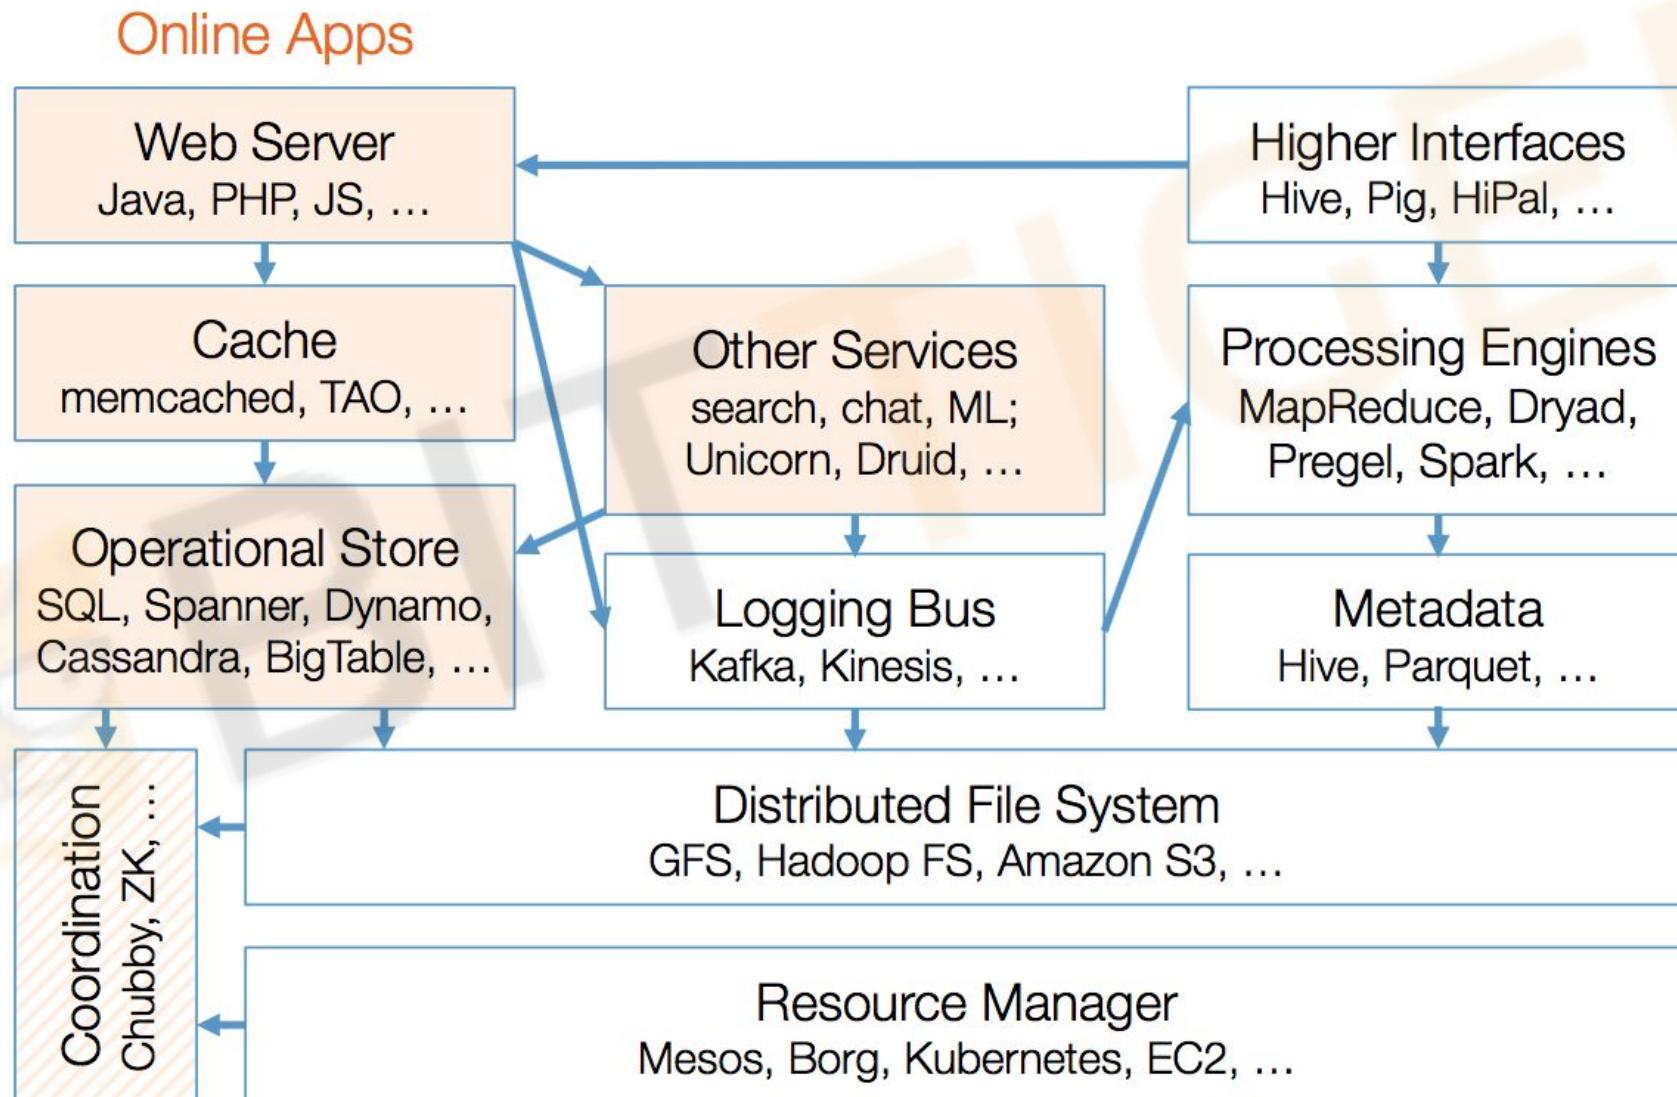
Scalable architecture



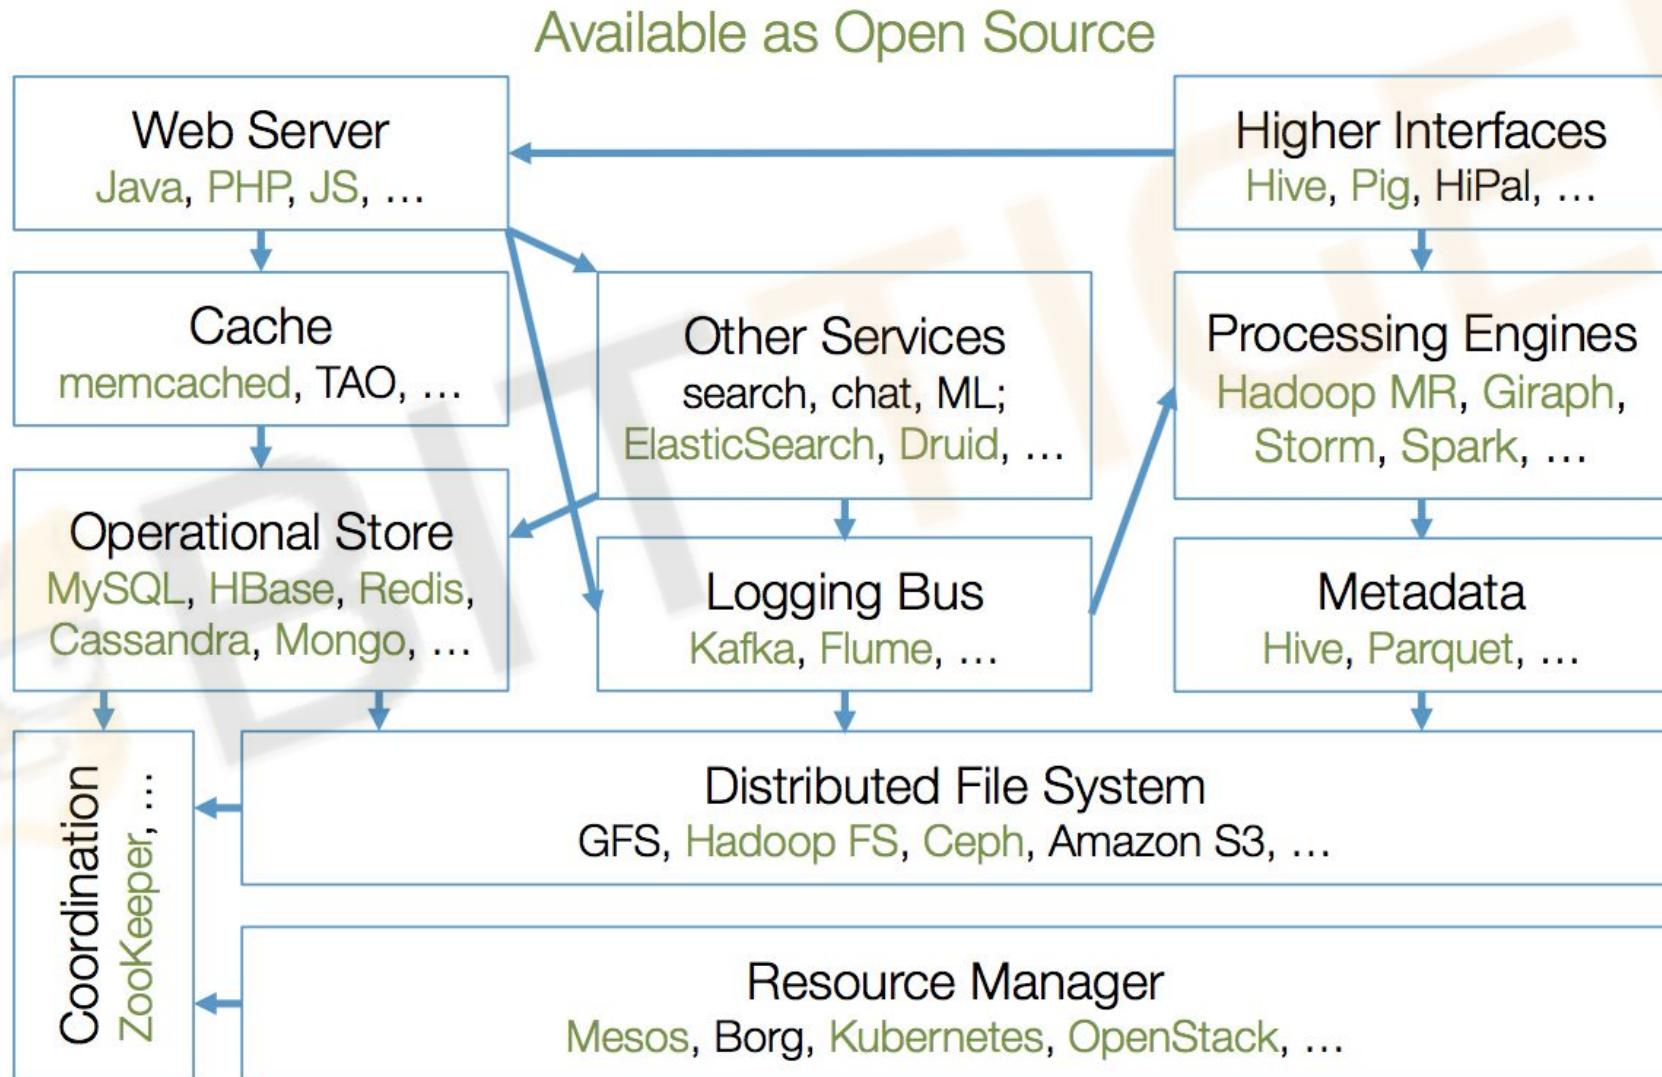
Scalable architecture



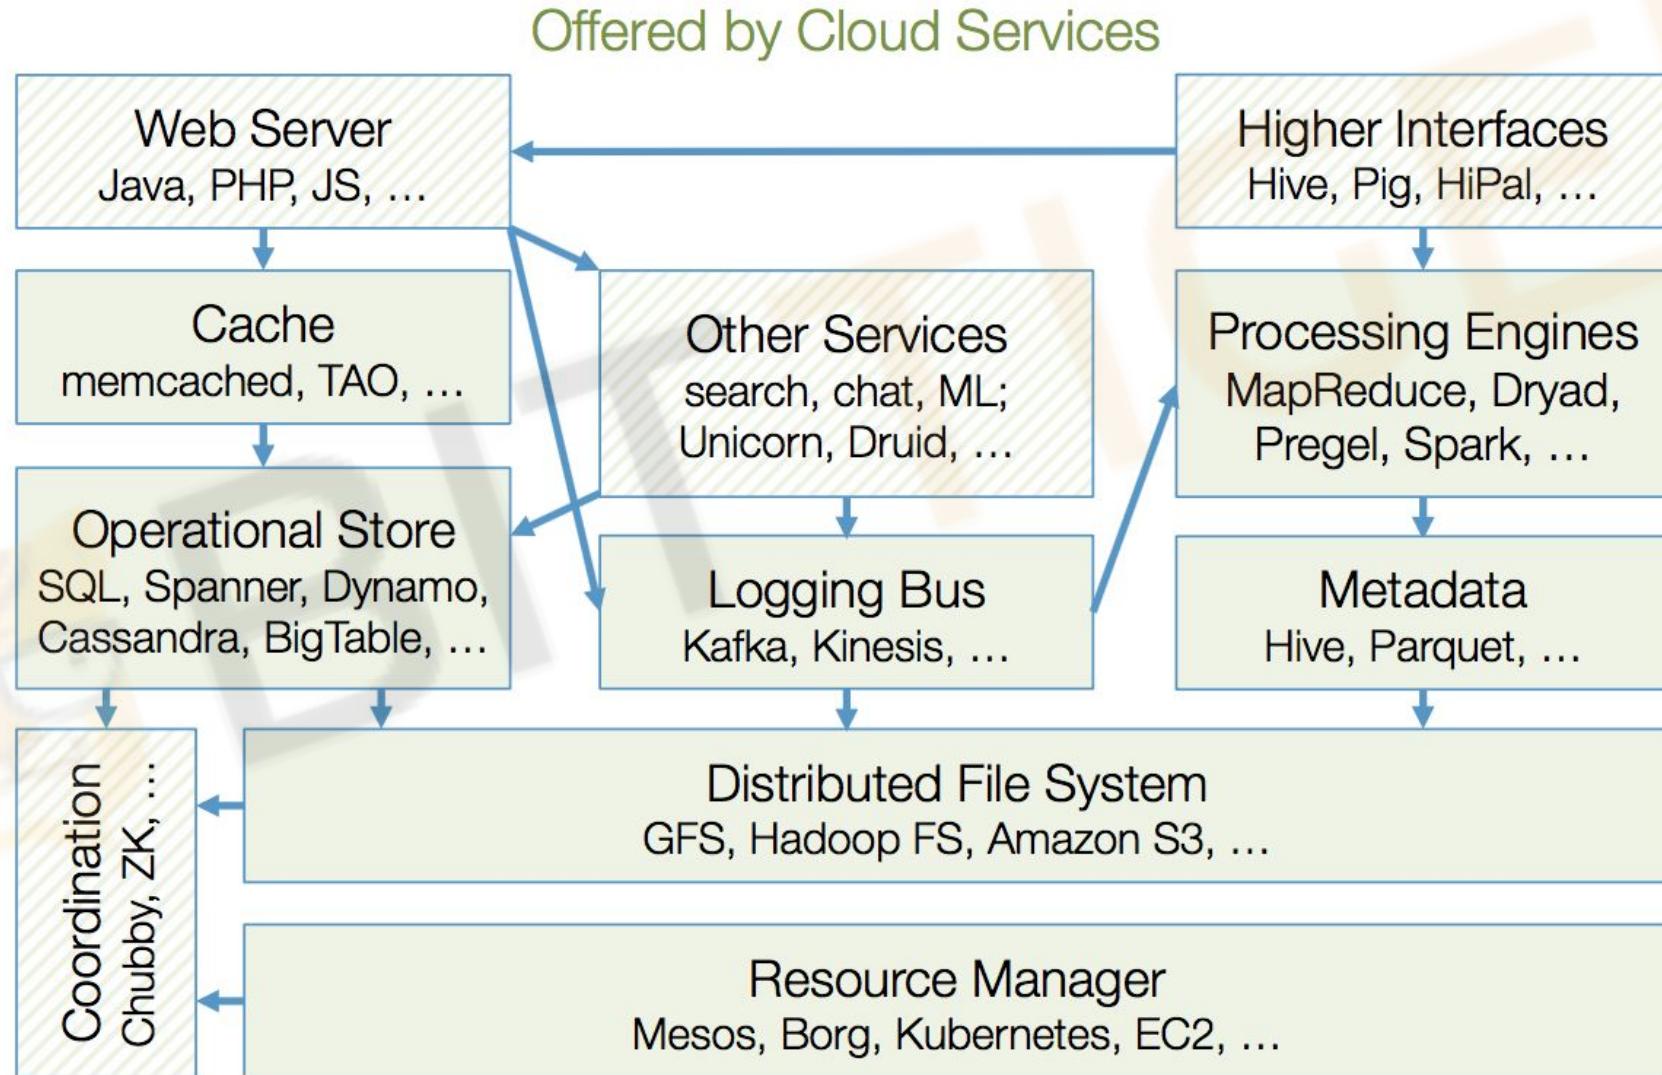
Scalable architecture



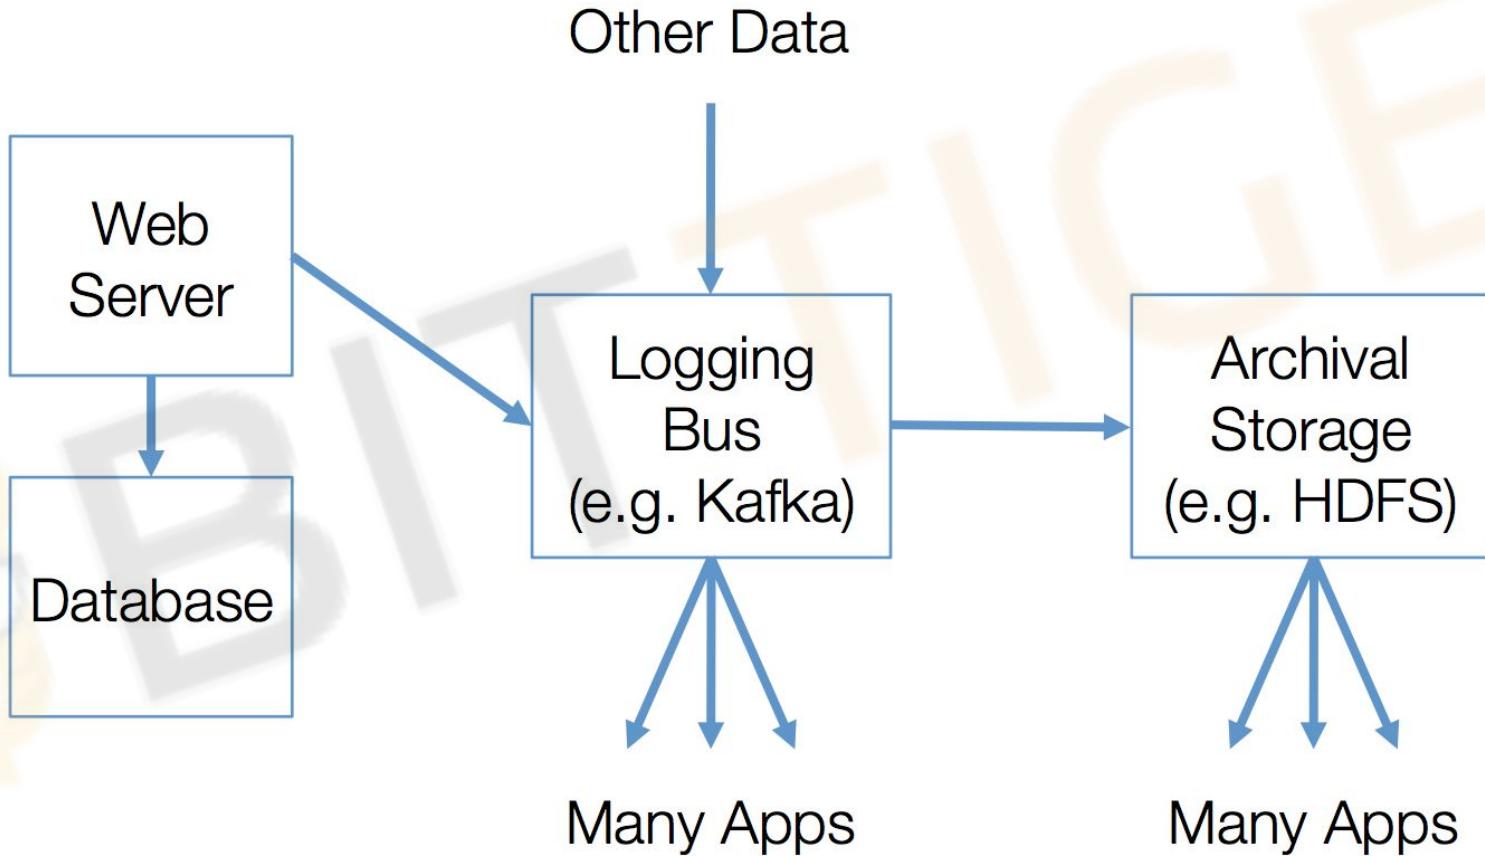
Scalable architecture



Software stack



Example architecture



Recommended reading: scaling facebook

- https://people.csail.mit.edu/matei/courses/2015/6.S897/readings/facebook_hanca2012.pdf
- Challenges of scaling facebook
 - Data is connected
 - Complex infrastructure
- Scaling facebook
 - Web: HipHop VM Services: Pull approach
 - Cache: Tao, Memcached
 - Storage: Haystack, f4

Delayed task scheduler

Delayed task scheduler

```
public interface Scheduler{  
    void schedule(Task t, long delayMs);  
}  
  
public interface Task{  
    void run();  
}
```

Delayed task scheduler: attempt 1

```
public class SchedulerImpl implements Scheduler{  
    public void schedule(Task t, long delayMs) {  
        try{  
            Thread.sleep(delayMs);  
            t.run();  
        }catch(InterruptedException e) {  
            //ignore  
        }  
    }  
}
```

Delayed task scheduler: attempt 1

```
public class SchedulerImpl implements Scheduler{  
    public void schedule(Task t, long delayMs) {  
        try{  
            Thread.sleep(delayMs);  
            t.run();  
        }catch(InterruptedException e){  
            //ignore  
        }  
    }  
}
```

Sleep for delayMs, then execute the task

Delayed task scheduler : attempt 1

```
public class TaskImpl implements Task{  
    private int id;  
    public TaskImpl(int id){  
        this.id = id;  
    }  
    public void run(){  
        System.out.println("Task " + id + " is running!");  
    }  
}
```

Delayed task scheduler : attempt 1

```
public class SchedulerImpl implements Scheduler{  
    public void schedule(Task t, long delayMs) {...}  
    public static void main(String[] args){  
        Scheduler scheduler = new SchedulerImpl();  
        Task t1 = new TaskImpl(1);  
        Task t2 = new TaskImpl(2);  
        scheduler.schedule(t1,10000);  
        scheduler.schedule(t2,1);  
    }  
}
```

Delayed task scheduler: attempt 1

```
public class SchedulerImpl implements Scheduler{  
    public void schedule(Task t, long delayMs) {...}  
    public static void main(String[] args){  
        Scheduler scheduler = new SchedulerImpl();  
        Task t1 = new TaskImpl(1);  
        Task t2 = new TaskImpl(2);  
        scheduler.schedule(t1,10000);  
        scheduler.schedule(t2,1);  
    }  
}
```

Main thread in Timedwaiting state for 1000ms

Delayed task scheduler : attempt 1

```
public class SchedulerImpl implements Scheduler{  
    public void schedule(Task t, long delayMs) {...}  
    public static void main(String[] args) {  
        Scheduler scheduler = new SchedulerImpl();  
        Task t1 = new TaskImpl(1);  
        Task t2 = new TaskImpl(2);  
        scheduler.schedule(t1,10000);  
        scheduler.schedule(t2,1);  
    }  
}
```

Task t2 will be delayed 10001ms

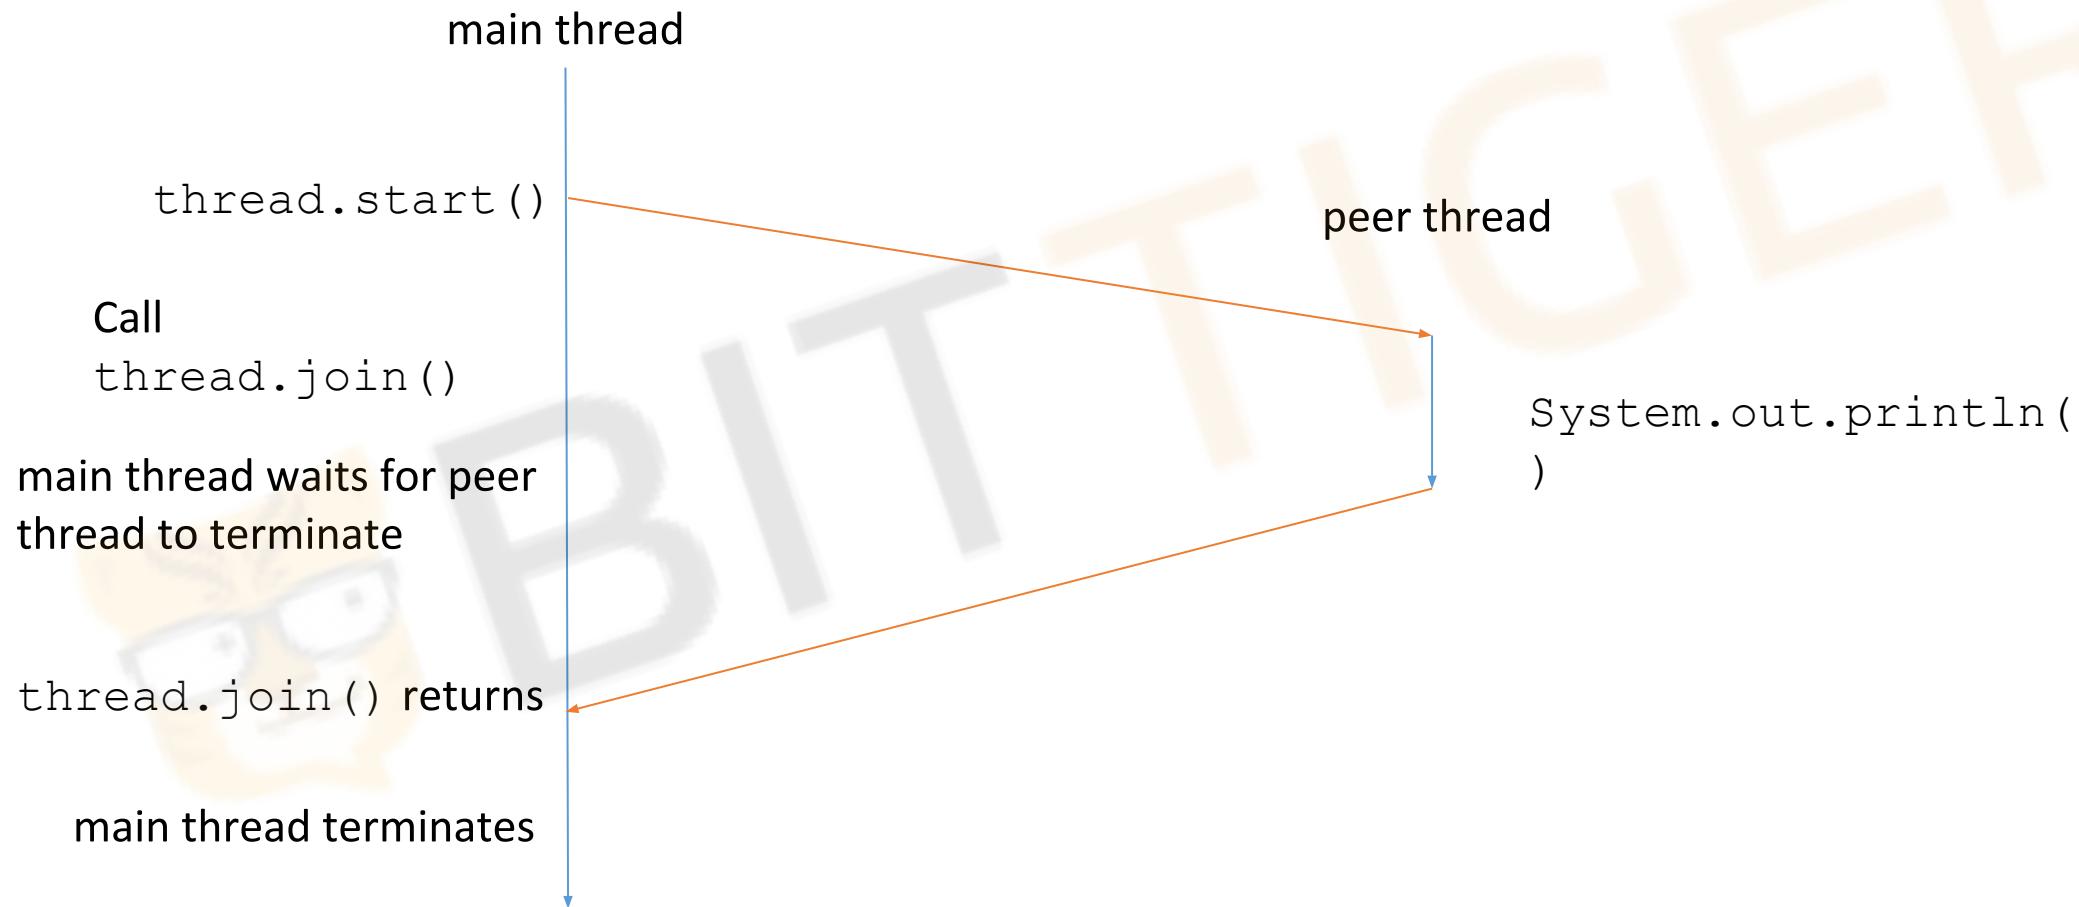
Delayed task scheduler : attempt 1

- Main thread is in Timedwaiting state for delayMs for each call of **schedule ()**.
- Only one thread, very low CPU utilization.
- Also, this is not working as later call
- How about sleeping in other threads?

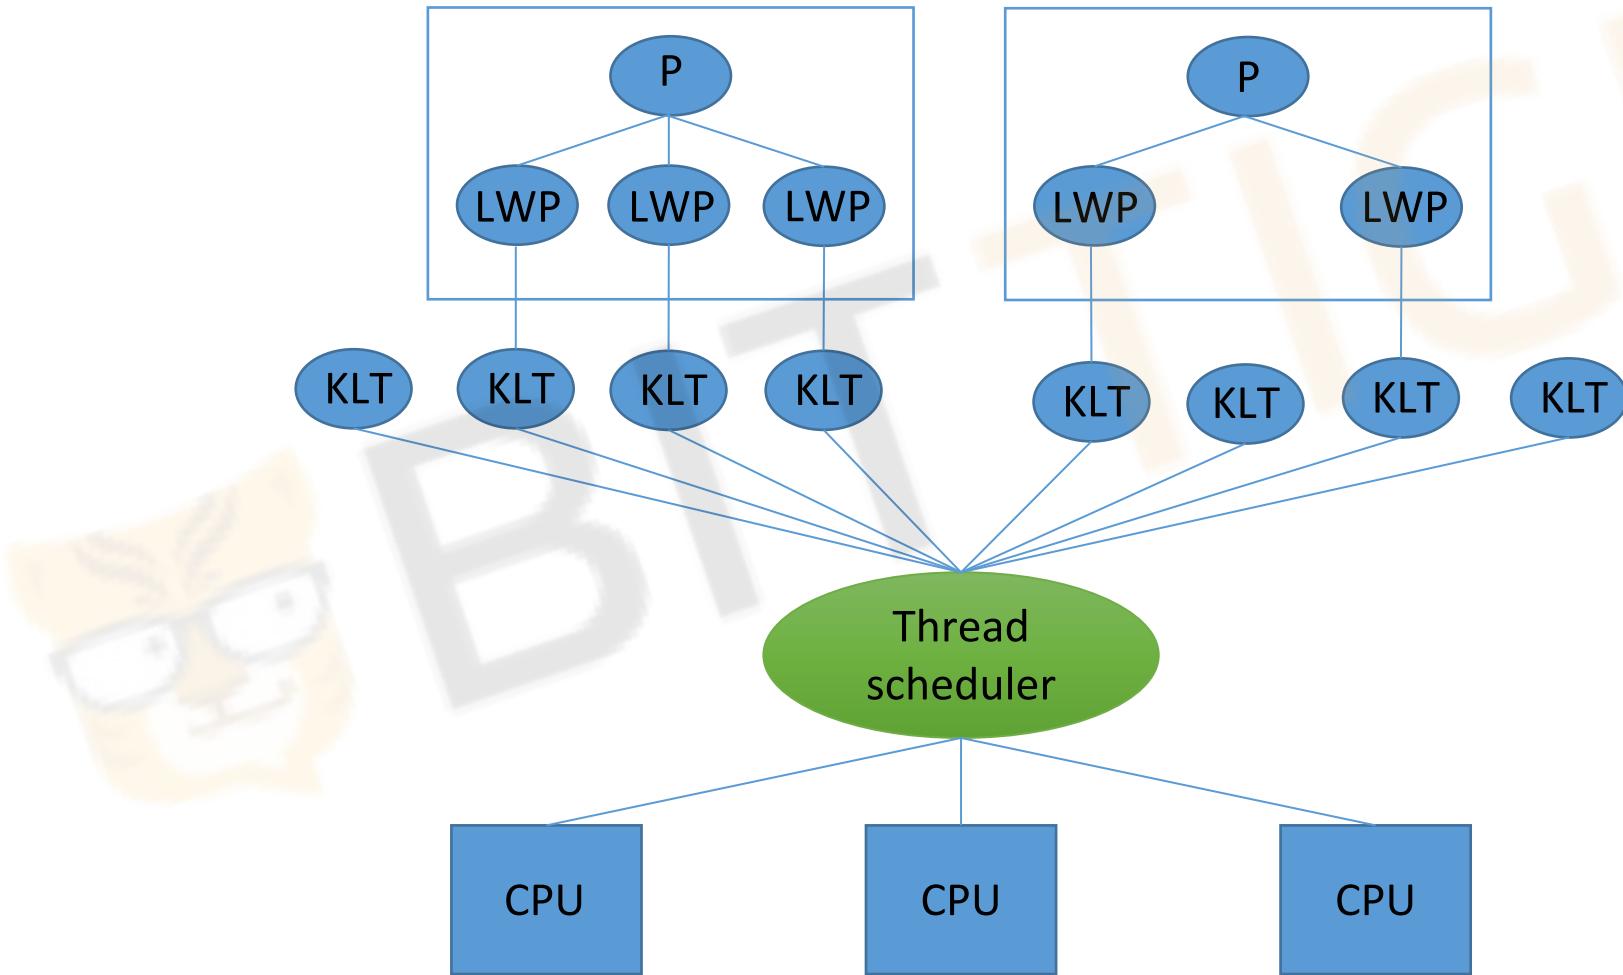
Java Threads

- Java Threads
 - Thread class
 - Runnable interface
- Thread creation
 - Override Thread class
 - Implement Runnable interface
- Hello World

Execution of threaded “Hello world”



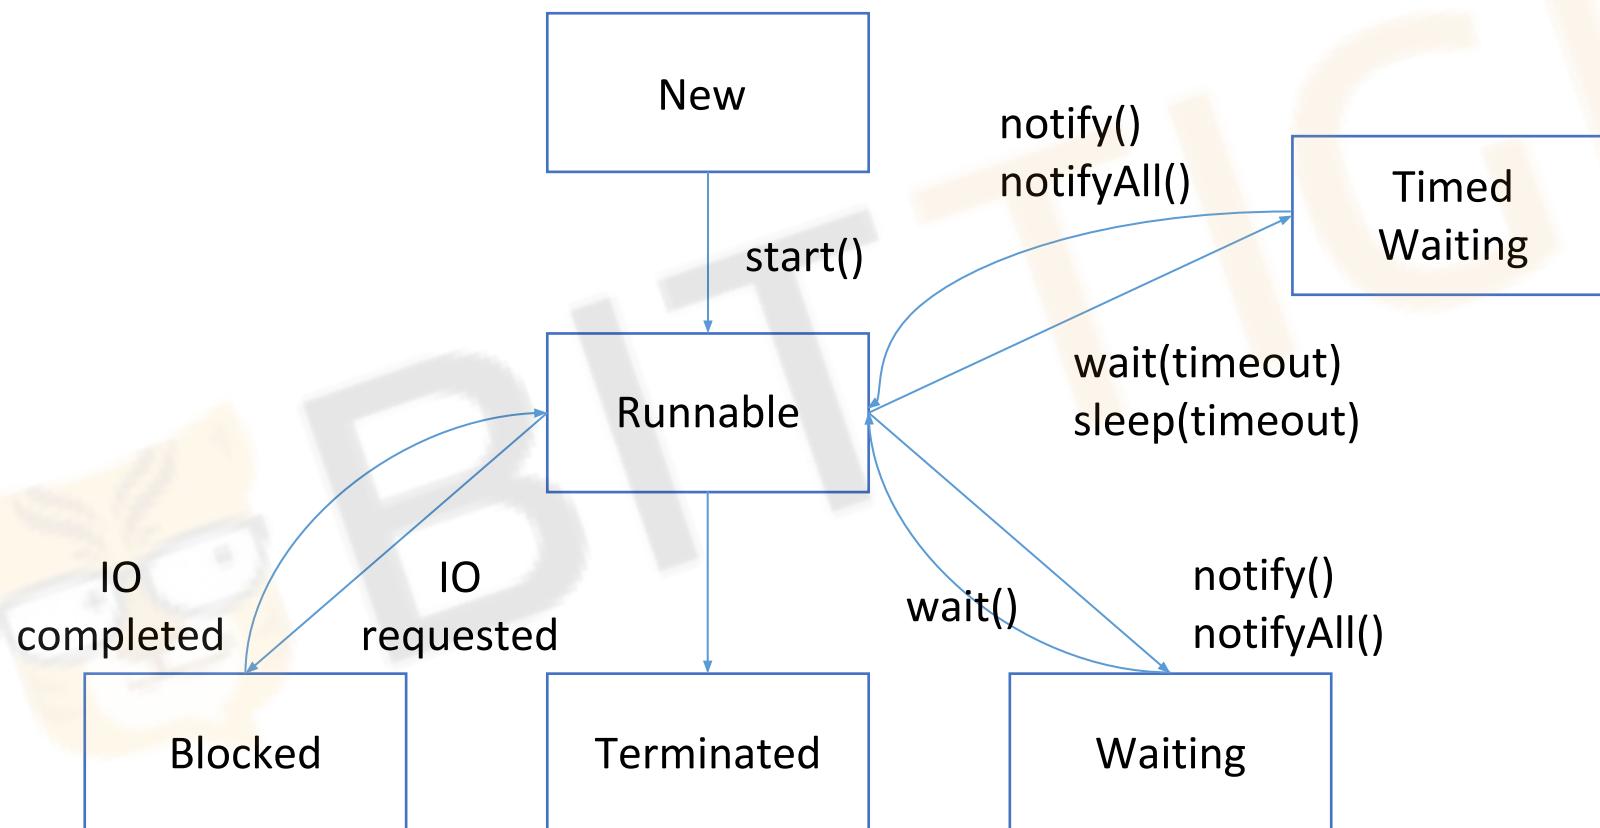
Threads implementation



Java thread states

- Java threads have 6 states, at any point of time, each thread can only be in one of the states
 - New
 - Runnable
 - Waiting
 - Timed waiting
 - Blocked
 - Terminated

Java thread lifecycle



Java thread lifecycle

```
Thread t = new Thread(); // New  
t.start(); // Runnable  
//Become running if scheduled by OS  
Thread.sleep(1000); // Timed waiting  
wait(); //Waiting  
System.out.println("Hello world"); //Blocked
```

Java threads summary

- Thread class and Runnable interface.
- Can use anonymous class to simplify thread creation.
- Java thread states: New, Runnable, Blocked, Waiting, Timed waiting, Terminated.
- A thread is in blocked state when performing IO. When blocked, the thread gives the CPU to other threads.
- In Linux, Java threads are mapped to kernel level threads (KLT) and scheduled by OS.

Synchronization

- Safety
 - Nothing bad happens ever
- Liveness
 - Something good happens eventually



BIT TIGER

Mutex

- Synchronized method

```
public class A {  
    public synchronized void foo() { ... }  
    public synchronized void bar() { ... }  
    public void foobar() { ... }  
}  
A a1 = new a1();  
T1: a1.foo();  
T2: a1.bar(); // T2 is blocked  
T3: a1.foobar(); T3 is not blocked
```

Mutex

- Synchronized code block
 - `synchronized(this) {};`
 - `Object lock = new Object(); synchronized(lock) {...};`
 - `class A {}; synchronized(A.class) {...};`
- Only one thread can access the synchronized block, other threads will be blocked.

Mutex

- Synchronized method/block is reentrant to the same thread.
- Other threads are blocked.
- Java threads maps to kernel threads, so the kernel is in charge of thread state transition.
- user mode -> kernel mode, so synchronized is heavy weight lock.
- Changes in a synchronized method/block are made visible to other threads when exiting the synchronized block.

Mutex

- class A {
 public void foo();
}
- A a1 = new A();
- A a2 = new A();
- a1.foo(); // T1
- a2.foo(); // T2 blocked
- We want to make sure that that only one thread can access foo() all
all instances of A.

Reentrant lock

- `synchronized` is equivalent to reentrant lock

```
Lock lock = new ReentrantLock();  
try {  
    lock.lock();  
    do something  
} finally {  
    lock.unlock();  
}
```

Condition variable

- A condition variable is an explicit queue that threads can put themselves on when some state of execution (i.e., some condition) is not as desired (by waiting on the condition).
- some other thread, when it changes said state, can then wake one (or more) of those waiting threads and thus allow them to continue.
- `wait()`
- `wait(long timeout)`
- `notify()`
- `notifyAll()`

Condition variable

- Conditional variable should be used inside a synchronized block.
- In Java, each object is associated with one and only one condition variable.
- This is a restriction, you may need to use explicit locking to get around this.

Condition variable

- Multiple condition variables can be associated with a reentrant lock.
- This is useful for implementing blocking queue.
- More about this in the next lecture.

Volatile variables

- Volatile semantics
 - Changes to volatile variables are visible to other threads immediately.
 - Volatile eliminate instruction reordering.
- Volatile use cases
 - Single writer.
 - The result doesn't depend on the current value.

Atomic variables

- For an atomic operation, all threads either see the value of before assignment, or the value after assignment.
- Compare and swap (CAS)

Atomic variables

- ```
public final int incrementAndGet() {
 for(;;) {
 int current = get();
 int next = current + 1;
 if (compareAndSet(current, next))
 return next;
 }
}
```
- ABA problem
  - A variable is value A when first read, and stays at A when assigning.
  - A -> B -> A

# Delayed task scheduler : attempt 2

```
public class SchedulerImpl implements Scheduler{
 public void schedule(Task t, long delayMs) {
 Thread t = new Thread(new Runnable() {
 public void run() {
 try{
 Thread.sleep(delayMs);
 t.run();
 }catch(InterruptedException e) {
 //ignore
 }
 }
);
 t.start();
 }
}
```

# Delayed task scheduler : attempt 2

```
public class SchedulerImpl implements Scheduler{
 public void schedule(Task t, long delayMs) {
 Thread t = new Thread(new Runnable() {
 public void run() {
 try{
 Thread.sleep(delayMs);
 t.run();
 }catch(InterruptedException e){
 //ignore
 }
 }
 };
 t.start();
 }
}
```

**Create a new thread for each schedule call**

# Delayed task scheduler : attempt 2

- No blocking when calling schedule()
- What happens if we call schedule many times?
- A lot of thread creation overhead.
- Can be alleviated by using a thread pool, but still not ideal.

# Delayed task scheduler: attempt 3

- Use a priority queue and a background thread.



BIT-TIGER

# Task

```
class Task implements Comparable<Task>{
 private long timeToRun;
 private int id;
 public void run() {
 System.out.println("Running task " + id);
 }
 public int compareTo(Task other) {
 return (int) (timeToRun - other.getTimeToRun());
 }
 ...
}
```

# Task

```
class Task implements Comparable<Task>{
 private long timeToRun;
 private int id;
 public void run() {
 System.out.println("Running task " + id);
 }
 public int compareTo(Task other) {
 return (int) (timeToRun - other.getTimeToRun());
 }
 ...
}
```

The time to run the task

# Task

```
class Task implements Comparable<Task>{
 private long timeToRun;
 private int id;
 public void run() {
 System.out.println("Running task " + id);
 }
 public int compareTo(Task other) {
 return (int) (timeToRun - other.getTimeToRun());
 }
 ...
}
```

Task id

# Task

```
class Task implements Comparable<Task>{
 private long timeToRun;
 private int id;
 public void run() {
 System.out.println("Running task " + id);
 }
 public int compareTo(Task other) {
 return (int) (timeToRun - other.getTimeToRun());
 }
 ...
}
```

**Dummy run method**

# Task

```
class Task implements Comparable<Task>{
 private long timeToRun;
 private int id;
 public void run() {
 System.out.println("Running task " + id);
 }
 public int compareTo(Task other) {
 return (int) (timeToRun - other.getTimeToRun());
 }
 ...
}
```

Order tasks by time to run

# Scheduler: fields

```
public class Scheduler{
 private PriorityQueue<Task> tasks;
 private final Thread taskRunnerThread;
 private volatile boolean running;
 private final AtomicInteger taskId;
 public Scheduler(){...}
 public void schedule(Task task, long delayMs){...}
 public void stop(){...}
 private class TaskRunner implements Runnable{...}
}
```

# Scheduler: fields

```
public class Scheduler{
 private PriorityQueue<Task> tasks;
 private final Thread taskRunnerThread;
 private volatile boolean running;
 private final AtomicInteger taskId;
 public Scheduler(){...}
 public void schedule(Task task, long delayMs){...}
 public void stop(){...}
 private class TaskRunner implements Runnable{...}
}
```

**Priority queue to order tasks by time to run**

# Scheduler: fields

```
public class Scheduler{
 private PriorityQueue<Task> tasks;
 private final Thread taskRunnerThread;
 private volatile boolean running;
 private final AtomicInteger taskId;
 public Scheduler(){...}
 public void schedule(Task task, long delayMs){...}
 public void stop(){...}
 private class TaskRunner implements Runnable{...}
}
```

**Background thread to run tasks**

# Scheduler: fields

```
public class Scheduler{
 private PriorityQueue<Task> tasks;
 private final Thread taskRunnerThread;
 private volatile boolean running;
 private final AtomicInteger taskId;
 public Scheduler(){...}
 public void schedule(Task task, long delayMs){...}
 public void stop(){...}
 private class TaskRunner implements Runnable{...}
}
```

**State indicating the scheduler is running**

# Scheduler: fields

```
public class Scheduler{
 private PriorityQueue<Task> tasks;
 private final Thread taskRunnerThread;
 private volatile boolean running;
 private final AtomicInteger taskId;
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

**Task id to assign to submitted tasks**

# Scheduler: constructor

```
public class Scheduler{
 public Scheduler(){
 tasks = new PriorityQueue<>();
 taskRunnerThread = new Thread(new TaskRunner());
 running = true;
 taskId = new AtomicInteger(0);
 taskRunnerThread.start();
 }
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

# Scheduler: constructor

```
public class Scheduler{
 public Scheduler(){
 tasks = new PriorityQueue<>();
 taskRunnerThread = new Thread(new TaskRunner());
 running = true;
 taskId = new AtomicInteger(0);
 taskRunnerThread.start();
 }
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

Create the background task runner

# Scheduler: constructor

```
public class Scheduler{
 public Scheduler(){
 tasks = new PriorityQueue<>();
 taskRunnerThread = new Thread(new TaskRunner());
 running = true;
 taskId = new AtomicInteger(0);
 taskRunnerThread.start();
 }
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

**Set running to true**

# Scheduler: constructor

```
public class Scheduler{
 public Scheduler(){
 tasks = new PriorityQueue<>();
 taskRunnerThread = new Thread(new TaskRunner());
 running = true;
 taskId = new AtomicInteger(0);
 taskRunnerThread.start();
 }
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

Initialize task id

# Scheduler: constructor

```
public class Scheduler{
 public Scheduler(){
 tasks = new PriorityQueue<>();
 taskRunnerThread = new Thread(new TaskRunner());
 running = true;
 taskId = new AtomicInteger(0);
 taskRunnerThread.start();
 }
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

**Start task runner thread**

# Scheduler: schedule()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {
 long timeToRun = System.currentTimeMillis() + delayMs;
 task.setTimeToRun(timeToRun);
 task.setId(taskId.incrementAndGet());
 synchronized(this) {
 tasks.offer(task);
 this.notify();
 }
 }
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

# Scheduler: schedule()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {
 long timeToRun = System.currentTimeMillis() + delayMs;
 task.setTimeToRun(timeToRun);
 task.setId(taskId.incrementAndGet());
 synchronized(this) {
 tasks.offer(task);
 this.notify();
 }
 }
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

**Set time to run and assign task id**

# Scheduler: schedule()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {
 long timeToRun = System.currentTimeMillis() + delayMs;
 task.setTimeToRun(timeToRun);
 task.setId(taskId.incrementAndGet());
 synchronized(this) {
 tasks.offer(task);
 this.notify();
 }
 }
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

Put the task in queue

# Scheduler: stop()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {...}
 public void stop(){
 synchronized(this) {
 running = false;
 this.notify();
 }
 taskRunnerThread.join();
 }
 private class TaskRunner implements Runnable{...}
}
```

# Scheduler: stop()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {...}
 public void stop(){
 synchronized(this){
 running = false;
 this.notify();
 }
 taskRunnerThread.join();
 }
 private class TaskRunner implements Runnable{...}
}
```

**Notify the task runner as it may be in wait()**

# Scheduler: stop()

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {...}
 public void stop(){
 synchronized(this){
 running = false;
 this.notify();
 }
 taskRunnerThread.join();
 }
 private class TaskRunner implements Runnable{...}
}
```

**Wait for the task runner to terminate**

# Scheduler: TaskRunner

```
public class Scheduler{
 public Scheduler() {...}
 public void schedule(Task task, long delayMs) {...}
 public void stop() {...}
 private class TaskRunner implements Runnable{...}
}
```

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run() {
 while(running) {
 synchronized(Scheduler.this) {
 while(running&&tasks.isEmpty()) {Scheduler.this.wait(); }
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(t.getTimeToRun() < now) {
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() -now);
 }
 }
 }
 }
}
```

# TaskRunner

```
private class TaskRunner implements Runnable {
 public void run() {
 while(running) {
 synchronized(Scheduler.this) {
 while(running&&tasks.isEmpty()) {Scheduler.this.wait(); }
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(t.getTimeToRun() < now) {
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun () -now);
 }
 }
 }
 }
}
```

**Inner class of Scheduler, implements Runnable**

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run(){
 while(running){
 synchronized(Scheduler.this){
 while(running&&tasks.isEmpty()){Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now){
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() - now);
 }
 }
 }
 }
}
```

The task runner keeps running until stopped from main thread

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run(){
 while(running){
 synchronized(Scheduler.this){
 while(running&&tasks.isEmpty()){Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now){
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() - now);
 }
 }
 }
 }
}
```

Need to synchronize with main thread

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run(){
 while(running){
 synchronized(Scheduler.this){
 while(running&&tasks.isEmpty()) {Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now) {
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() - now);
 }
 }
 }
 }
}
```

**Task runner is blocked when no tasks in queue**

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run(){
 while(running){
 synchronized(Scheduler.this){
 while(running&&tasks.isEmpty()){Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now){
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() - now);
 }
 }
 }
 }
}
```

Check the first task in queue

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run() {
 while(running) {
 synchronized(Scheduler.this) {
 while(running&&tasks.isEmpty()) {Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now) {
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun()-now);
 }
 }
 }
 }
}
```

**Delay exhausted, execute task**

# TaskRunner

```
private class TaskRunner implements Runnable{
 public void run(){
 while(running){
 synchronized(Scheduler.this){
 while(running&&tasks.isEmpty()){Scheduler.this.wait();}
 long now = System.currentTimeMillis();
 Task t = tasks.peek();
 if(running&&t.getTimeToRun() < now){
 tasks.poll();
 t.run();
 }else{
 Scheduler.this.wait(t.getTimeToRun() - now);
 }
 }
 }
 }
}
```

No task executable, wait

# The wait(timeout) method

```
q.wait(timeout);
```

- Releases lock on q
- Sleeps (gives up processor)
- Awakens (resumes running) by notify() and notifyAll()
- Awakens after timeout
- Reacquires lock & returns

# Why wait(timeout)?

```
Scheduler scheduler = new Scheduler();
scheduler.schedule(new Task(),1000000);
scheduler.schedule(new Task(),1000);
```



BIT

# Delayed task scheduler summary

- Single threaded implementation is not working because main the thread is blocked for each **schedule()** call.
- Creating a new thread for each **schedule()** call is not scalable with system resources.
- Use a priority queue and background task runner thread.
- sleep(timeout) vs wait(timeout)